

# SQL – Data Definition Language

---

Carla Teixeira Lopes

Bases de Dados

Mestrado Integrado em Engenharia Informática e Computação, FEUP

# Agenda

---

SQL Introduction

Defining a Relation Schema in SQL

Constraints

# (Integrity) Constraints

---

Part of the SQL standard but systems vary considerable

Impose restrictions on allowable data, beyond those imposed by structure and types

## Examples

$0.0 < GPA \leq 4.0$

$enrolment < 50,0000$

Decision: 'Y', 'N' or NULL

$Major = 'CS' \Rightarrow decision = NULL$

$sizeHS < 200 \Rightarrow not\ admitted\ in\ colleges\ with\ enroll > 30,000$

# Why use constraints?

---

Data-entry errors (inserts)

Correctness criteria (updates)

Enforce consistency

Tell system about data – store, query processing

# Classification of constraints

---

Non-null constraints

Key constraints

Attribute-based and tuple-based constraints

Referential integrity (foreign key)

General assertions

# Declaring and enforcing constraints

---

## Declaration

- With original schema

  - Checked after bulk loading

- Or later

  - Checked at the time it's declared on the current state of the DB

## Enforcement

- Check after every modification

  - Check only the dangerous ones

  - If there is a constraint on one table, there is no need to check updates on another tables

- Deferred constraint checking

  - Instead of checking after every modification, checking is done after every transaction

# Non-null constraint

---

Defines that a column does not have NULL values

```
CREATE TABLE <table_name> (  
  <column_name> <data_type> NOT NULL,  
  ...  
  <column_name> <data_type>  
);
```

# Example

---

```
CREATE TABLE MovieStar (  
    id            INTEGER,  
    name          CHAR(30) NOT NULL,  
    address       VARCHAR(255),  
    gender        CHAR(1) DEFAULT '?',  
    birthdate     DATE  
);
```



# Primary key (PK) constraint

---

We can define one, and only one, primary key for a table

```
CREATE TABLE <table_name> (  
  <column_name> <data_type> PRIMARY KEY,  
  ...  
  <column_name> <data_type>  
);
```

A PK cannot be NULL and cannot have repeated values

In SQLite, INTEGER PK will auto increment if a NULL value is inserted in the PK column

# Example

---

```
CREATE TABLE MovieStar (  
    id          INTEGER PRIMARY KEY,  
    name        CHAR(30),  
    address     VARCHAR(255),  
    gender      CHAR(1) DEFAULT '?',  
    birthdate   DATE  
);
```

# Multiple Column PK

---

If a primary key is composed by more than one column

```
CREATE TABLE <table_name> (  
  <column_name> <data_type>,  
  ...  
  <column_name> <data_type>,  
  PRIMARY KEY (<column_name>, <column_name>)  
);
```

# Example

---

```
CREATE TABLE MovieStar (  
    name      CHAR(30),  
    address   VARCHAR(255),  
    gender    CHAR(1) DEFAULT '?',  
    birthdate DATE,  
    PRIMARY KEY (name, birthdate)  
);
```

# ROWID in SQLite

---

If a table is created without specifying the WITHOUT ROWID option, SQLite adds an implicit column called rowid that stores 64-bit signed integer

The rowid is a key that uniquely identifies the row in its table

Can be accessed using ROWID, \_ROWID\_, or OID (except if ordinary columns use these names)

On an INSERT, if the ROWID is not explicitly given a value, then it will be filled automatically with an unused integer greater than 0, usually one more than the largest ROWID currently in use.

# SQLite PK and ROWID

---

If a table has a one-column PK defined as INTEGER, this PK column becomes an alias for the rowid column

Tables with rowid are stored as a B-Tree using rowid as the key

- Retrieving and sorting by rowid are very fast

- Faster than using any other PK or indexed value

# SQLite Autoincrement

---

Imposes extra CPU, memory, disk space, and disk I/O overhead

If an AUTOINCREMENT keyword appears after INTEGER PRIMARY KEY, that changes the ROWID assignment algorithm to prevent the reuse of ROWIDs

The purpose of AUTOINCREMENT is to prevent the reuse of ROWIDs from previously deleted rows

Should be avoided if not strictly needed

# Unique key constraint

---

We can define multiple unique keys for a table

```
CREATE TABLE <table_name> (  
  <column_name> <data_type> UNIQUE,  
  ...  
  <column_name> <data_type>  
);
```

Unique Keys allow NULL values

The SQL standard and most DBMS do allow repeated NULL values in UNIQUE keys



# Example

---

```
CREATE TABLE MovieStar (  
    id          INTEGER PRIMARY KEY,  
    name        CHAR(30),  
    address     VARCHAR(255) UNIQUE,  
    gender      CHAR(1) DEFAULT '?',  
    birthdate   DATE,  
    phone       CHAR(16) UNIQUE  
);
```

Address cannot have repeated values

Phone cannot have repeated values

# Multiple Column Unique Key

---

If a unique key is composed by more than one column

```
CREATE TABLE <table_name> (  
  <column_name> <data_type>,  
  ...  
  <column_name> <data_type>,  
  UNIQUE (<column_name>, <column_name>)  
);
```

# Example

---

```
CREATE TABLE MovieStar (  
    id          INTEGER PRIMARY KEY,  
    name        CHAR(30),  
    address     VARCHAR(255),  
    gender      CHAR(1) DEFAULT '?',  
    birthdate   DATE,  
    UNIQUE (name, birthdate),  
    UNIQUE (name, address)  
);
```

# Attribute-based check constraint

---

Constrain the value of a particular attribute

```
CREATE TABLE <table_name> (  
  <column_name> <data_type> CHECK <check_expression>,  
  ...  
  <column_name> <data_type>  
);
```

Checked whenever we insert or update a tuple

# Example

---

```
CREATE TABLE Student (  
    sID      INTEGER PRIMARY KEY,  
    sName    TEXT,  
    GPA      REAL CHECK (GPA<=4.0 and GPA>0.0),  
    sizeHS    INTEGER CHECK (sizeHS < 5000),  
);
```

GPA's must be less than or equal to 4.0 and greater than zero

The size of the high school must be less than five thousand

# Tuple-based check constraint

---

Constrain relationships between different values in each tuple

```
CREATE TABLE <table_name> (  
  <column_name> <data_type>,  
  ...  
  <column_name> <data_type>,  
  CHECK (<check_expression>  
);
```

Checked whenever we insert or update a tuple

# Example

---

```
CREATE TABLE Apply (  
    sID          INTEGER,  
    cName        TEXT,  
    major        TEXT,  
    decision     TEXT,  
    PRIMARY KEY (sID, cName, major),  
    CHECK (decision='N' or cName <>'Stanford' or major <>'CS')  
);
```

Either the decision is null or the college name is not Stanford or the major is not CS



There are no people who have applied to Stanford and been admitted to CS at Stanford

# Example

---

```
CREATE TABLE person (  
    id          INTEGER PRIMARY KEY,  
    name        TEXT,  
    salary      TEXT,  
    taxes       TEXT,  
    CHECK(taxes<salary)  
);
```

Taxes have to be lower than salary



# Subqueries and check constraints

---

```
CREATE TABLE Student (sID INTEGER PRIMARY KEY, sName  
TEXT, GPA REAL, sizeHS INTEGER);
```

```
CREATE TABLE Apply (  
    sID          INTEGER,  
    cName        TEXT,  
    major        TEXT,  
    decision     TEXT,  
    PRIMARY KEY (SID, cName, major),  
    CHECK (sID in (select sID from Student))  
);
```

Syntactically valid in the SQL standard but no SQL systems supports subqueries and check constraints.

# Kahoot time!

---

Any doubts?

# Readings

---

Jeffrey Ullman, Jennifer Widom, A first course in Database Systems 3<sup>rd</sup> Edition

Section 2.3 – Defining a Relation Schema in SQL

Section 2.5 – Constraints on Relations

Section 7.1 – Keys and Foreign Keys

Section 7.2 – Constraints on Attributes and Tuples

Section 7.3 – Modification if Constraints

Section 7.4 - Assertions