

# Relational Algebra

---

Carla Teixeira Lopes

Bases de Dados

Mestrado Integrado em Engenharia Informática e Computação, FEUP

Based on Jennifer Widom slides

# Agenda

---

Introduction to Relational Algebra

Operators

Alternate notations

Extensions to Relational Algebra

# What is Algebra?

---

Mathematical system consisting of:

Operands - variables or values from which new values can be constructed.

Operators - symbols denoting procedures that construct new values from given values.

# What is Relational Algebra?

---

An algebra whose operands are relations or variables that represent relations.

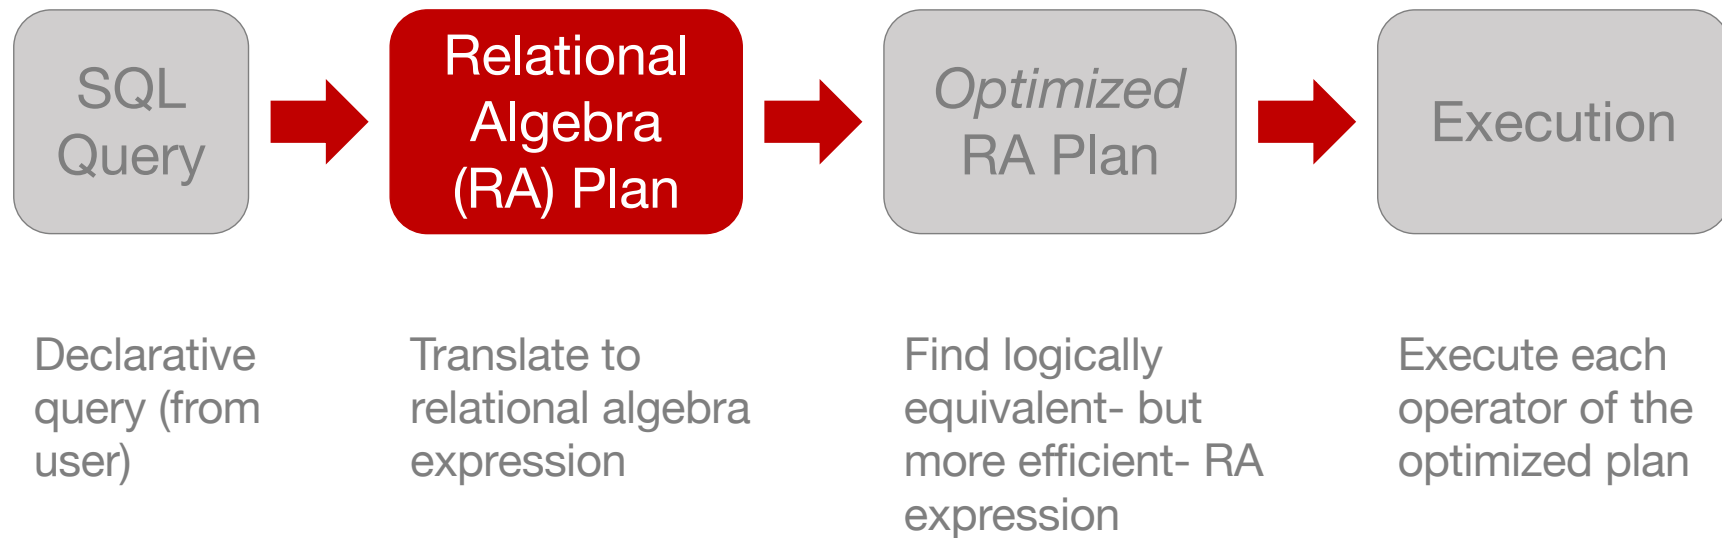
Operators are designed to do the most common things that we need to do with relations in a database.

The result is an algebra that can be used as a query language for relations.

# RDBMS Architecture

---

How does a SQL engine work?




# Relational Algebra

---

Formal language

Operates on relations and produce relations as a result

Operators are used to filter, slice and combine



# Agenda

---

Introduction to Relational Algebra

Operators

Alternate notations

Extensions to Relational Algebra

# College Admission Database

---

College (cName, state, enr)

Student (sID, sName, GPA, HS)

Apply (sID, cName, major, dec)

Demo in Relax: <https://dbis-uibk.github.io/relax/>

College

<u>cName</u>	state	enr

Student

<u>sID</u>	sName	GPA	HS

Apply


<u>sID</u>	<u>cName</u>	<u>major</u>	dec



# Simplest query: relation name

---

Student



Student

sID	sName	GPA	HS

## Select operator ( $\sigma$ )

---

Returns all tuples which satisfy a condition

Notation:  $\sigma_{condition}Relation$

The condition can involve =, <, ≤, >, ≥, <>

# Examples

---

Students with  $GPA > 3.7$

$\sigma_{GPA > 3.7}$  *Student*

Students with  $GPA > 3.7$  and  $HS < 1000$

$\sigma_{GPA > 3.7 \wedge HS < 1000}$  *Student*

Applications to Stanford CS major

$\sigma_{cName='Stanford' \wedge major='CS'}$  *Apply*

Student

sID	sName	GPA	HS
12	Mary	3.5	90
23	John	3.8	500
31	Jane	3.9	1000

Apply

sID	cName	major	dec
12	Stanford	CS	Y
23	MIT	CS	N
12	MIT	CS	N

# Project operator ( $\pi$ )

---

Picks certain columns

Notation:  $\pi_{A_1, \dots, A_n} \text{Relation}$

sID and decision of all applications

$\pi_{sID, dec} \text{Apply}$

Apply

sID	cName	major	dec
12	Stanford	CS	Y
23	MIT	CS	N
12	MIT	CS	N



sID	dec
12	Y
23	N
12	N

# Combining the Select and Project Operators

---

ID and name of students with GPA>3.7

$$\pi_{sID, sName} (\sigma_{GPA > 3.7} Student)$$

Redefinition of operators

$$\sigma_{condition}(Expression)$$

$$\pi_{A_1, \dots, A_n} (Expression)$$

Student

sID	sName	GPA	HS
12	Mary	3.5	90
23	John	3.8	50
31	Jane	3.9	1000

# Sets, Bags and Lists

---

## Sets

Only one occurrence of each element

Unordered elements

## Bags (or multisets)

More than one occurrence of an element

Unordered elements and their occurrences

## Lists

More than one occurrence of an element

Occurrences are ordered

# Duplicates

---

## Relational Algebra

Eliminates duplicates

Based on sets (although there is also a multiset relation algebra)

## SQL

Does not eliminate duplicates

Based on multisets or bags

## List of application majors and decisions

$\pi_{major,dec} Apply$

Apply

sID	cName	major	dec
12	Stanford	CS	Y
23	MIT	CS	N
12	MIT	CS	N

$\pi_{major,dec} Apply$

major	dec
CS	Y
CS	N

No duplicates



# Cross-product

---

Also known as Cartesian product

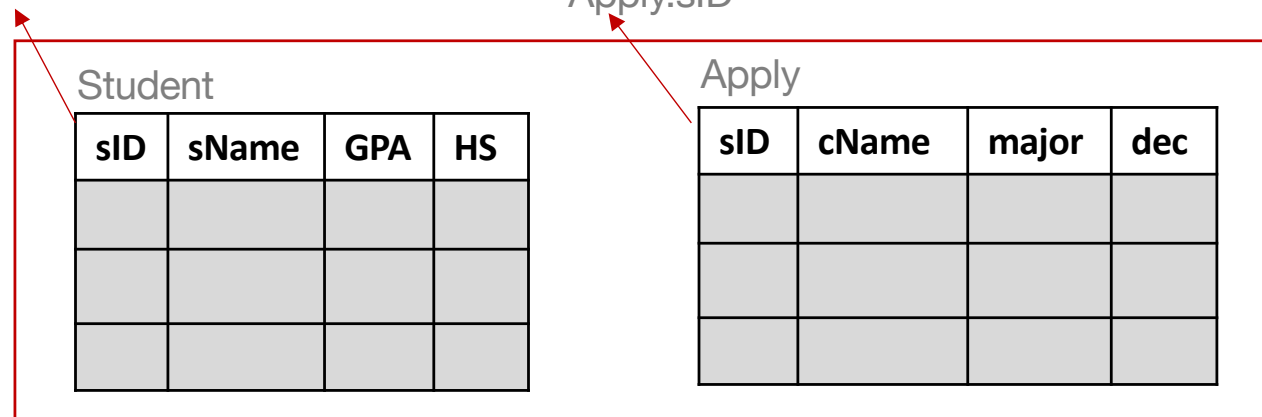
Notation: Rel1 x Rel2

Student x Apply

Attributes with the same name are prefaced with the name of the relation

Student.sID

Apply.sID





# Cross-product

---

One tuple for every combination of tuples from the student and apply relations

Student

sID	sName	GPA	HS

S tuples

Student x Apply



S x A tuples

Apply

sID	cName	major	dec

A tuples

# Example 1

---

Student

sID	sName	GPA	HS
12	Mary	3.5	90
23	John	3.8	50

Apply

sID	cName	major	dec
12	Stanford	CS	Y
23	MIT	CS	N



Student x Apply

Student.sID	sName	GPA	HS	Apply.sID	cName	major	dec
12	Mary	3.5	90	12	Stanford	CS	Y
12	Mary	3.5	90	23	MIT	CS	N
23	John	3.8	50	12	Stanford	CS	Y
23	John	3.8	50	23	MIT	CS	N

## Example 2

---

Names and GPAs of students with HS>100 who applied to CS and were rejected

Student x Apply

All combinations

$\sigma_{Student.sID=Apply.sID}(Student \times Apply)$

Combinations that make sense

$\sigma_{Student.sID=Apply.sID \wedge HS>100 \wedge major='CS' \wedge dec='N'}(Student \times Apply)$  Additional filtering

$\pi_{sName,GPA}(\sigma_{Student.sID=Apply.sID \wedge HS>100 \wedge major='CS' \wedge dec='N'}(Student \times Apply))$

Student

sID	sName	GPA	HS
12	Mary	3.5	90
23	John	3.8	5000

Apply

sID	cName	major	dec
12	Stanford	CS	Y
23	MIT	CS	N

# Natural Join

---

Operator: ⋈

Cross product enforcing equality on all attributes with same name

Eliminate one copy of duplicate attributes

College

cName	state	enr

Student

sID	sName	GPA	HS

Apply

sID	cName	major	dec

# Example 1

---

Student

sID	sName	GPA	HS
12	Mary	3.5	90
23	John	3.8	50

Apply

sID	cName	major	dec
12	Stanford	CS	Y
23	MIT	CS	N



*Student ⋈ Apply*

sID	sName	GPA	HS	cName	major	dec
12	Mary	3.5	90	Stanford	CS	Y
23	John	3.8	50	MIT	CS	N

## Example 2

---

Names and GPAs of students with  $HS > 100$  who applied to CS and were rejected

$Student \bowtie Apply$

$\sigma_{HS > 100 \wedge major = 'CS' \wedge dec = 'N'}(Student \bowtie Apply)$

$\pi_{sName, GPA}(\sigma_{HS > 100 \wedge major = 'CS' \wedge dec = 'N'}(Student \bowtie Apply))$

Student

sID	sName	GPA	HS
12	Mary	3.5	90
23	John	3.8	5000

Apply

sID	cName	major	dec
12	Stanford	CS	Y
23	MIT	CS	N

# Example 3

---

Names and GPAs of students with  $HS > 100$  who applied to CS at college with  $enr > 10,000$  and were rejected

$Student \bowtie (Apply \bowtie College)$

$\sigma_{HS > 100 \wedge major = 'CS' \wedge dec = 'N' \wedge enr > 10,000}(Student \bowtie (Apply \bowtie College))$

$\pi_{sName, GPA}(\sigma_{HS > 100 \wedge major = 'CS' \wedge dec = 'N' \wedge enr > 10,000}(Student \bowtie (Apply \bowtie College)))$

College

cName	state	enr
MIT	NULL	30000
Stanford	NULL	20000

Student

sID	sName	GPA	HS
12	Mary	3.5	90
23	John	3.8	5000

Apply

sID	cName	major	dec
12	Stanford	CS	Y
23	MIT	CS	N

# Natural Join

---

Given  $R(A, B, C, D)$ ,  $S(A, C, E)$ , what is the schema of  $R \bowtie S$  ?

Given  $R(A, B, C)$ ,  $S(D, E)$ , what is  $R \bowtie S$  ?

Given  $R(A, B)$ ,  $S(A, B)$ , what is  $R \bowtie S$  ?



# Natural Join does not add expressive power

---

Can be rewritten using the cross-product

$$Exp1 \bowtie Exp2 \equiv \pi_{schema(E1) \cup schema(E2)}(\sigma_{E1.A1=E2.A1 \wedge E1.A2=E2.A2 \wedge \dots}(Exp1 \times Exp2))$$

It is convenient in terms of notation

# Theta Join

---

A join that involves a predicate

Notation:  $\bowtie_{\theta}$

$$Exp_1 \bowtie_{\theta} Exp_2 \equiv \sigma_{\theta}(Exp_1 \times Exp_2)$$

Basic operation implemented in DBMS

Term “join” often means theta join

$\theta$  can be any condition

If  $\theta$  is an equality, the join is called an equi-join

# Example

---

Student

ID	sName	GPA	HS
12	Mary	3.5	90
23	John	3.8	50

Apply

sID	cName	major	dec
12	Stanford	CS	Y
23	MIT	CS	N



*Student* ⋈<sub>ID=sID</sub> *Apply*

ID	sName	GPA	HS	sID	cName	major	dec
12	Mary	3.5	90	12	Stanford	CS	Y
23	John	3.8	50	23	MIT	CS	N

# Semijoin

---

Notation:  $\bowtie$

$$Exp_1 \bowtie Exp_2 \equiv \pi_{A_1, \dots, A_n} (Exp_1 \bowtie Exp_2)$$

Where  $A_1, \dots, A_n$  are attributes in  $Exp_1$

Returns the tuples of  $Exp_1$  with a pair in  $Exp_2$

# Example

---

Student

sID	sName	GPA	HS
12	Mary	3.5	90
23	John	3.8	50
35	Jane	3.9	60

Apply

sID	cName	major	dec
12	Stanford	CS	Y
23	MIT	CS	N



*Student* ⋈ *Apply*

sID	sName	GPA	HS
12	Mary	3.5	90
23	John	3.8	50

# Union operator

---

Operator:  $\cup$

List of college and student names

Can we do it using previous operators?

$$\pi_{cName} College \cup \pi_{sName} Student$$

Combines information vertically

Technically, the two operands have to have the same schema

Not the case in the example above, but we'll correct it later

College

cName	state	enr
MIT	NULL	NULL
Washington	NULL	NULL

Student

sID	sName	GPA	HS
12	Mary	3.5	90
23	Washington	3.8	50

# Example


---

Student

sID	sName	GPA	HS
12	Mary	3.5	90
23	Washington	3.8	50


College

cName	state	enr
MIT	NULL	NULL
Washington	NULL	NULL



$\pi_{cName} College \cup \pi_{sName} Student$

cName
Mary
Washington
MIT



# Difference operator

---

Operator:  $-$


IDs of students who didn't apply anywhere

Student

sID	sName	GPA	HS
12	Mary	3.5	90
23	Washington	3.8	50


Apply

sID	cName	major	dec
12	Stanford	CS	Y



$$\pi_{sID} Student - \pi_{sID} Apply$$

sID
23






# Example

---

Names of students who didn't apply anywhere


$$\pi_{sName} Student - \pi_{sID} Apply ?$$

$$\pi_{sName} ((\pi_{sID} Student - \pi_{sID} Apply) \bowtie Student)$$


Schema equal to the student relation

Student

sID	sName	GPA	HS
12	Mary	3.5	90
23	Washington	3.8	50

$$\pi_{sID} Student - \pi_{sID} Apply$$


Apply

sID	cName	major	dec
12	Stanford	CS	Y

sID
23

# Intersection operator

---

Operator:  $\cap$

Names that are both a college name and a student name

$$\pi_{cName} College \cap \pi_{sName} Student$$

Technically, the two operands have to have the same schema

Not the case in the example above, but we'll correct it later

College

cName	state	enr
MIT	NULL	NULL
Washington	NULL	NULL

Student

sID	sName	GPA	HS
12	Mary	3.5	90
23	Washington	3.8	50

# Example


---

Student

sID	sName	GPA	HS
12	Mary	3.5	90
23	Washington	3.8	50


College

cName	state	enr
MIT	NULL	NULL
Washington	NULL	NULL



$\pi_{cName} College \cap \pi_{sName} Student$




cName
Washington



# Intersection doesn't add expressive power

---

$$E_1 \cap E_2 \equiv \textcolor{red}{E}_1 - (\textcolor{green}{E}_1 - \textcolor{green}{E}_2)$$



# Intersection doesn't add expressive power

---

$$E_1 \cap E_2 \equiv E_1 \bowtie E_2$$



Identical schema

Nevertheless, the intersection can be very useful in queries

# Kahoot time!

---

Any doubts?

# Readings

---

Jeffrey Ullman, Jennifer Widom, A first course in  
Database Systems 3<sup>rd</sup> Edition

Section 2.4 – An Algebraic Query Language

Section 5.2 – Extended Operators of Relational Algebra