

# SQL – Data Manipulation Language

---

Carla Teixeira Lopes

Bases de Dados

Mestrado Integrado em Engenharia Informática e Computação, FEUP

Based on Jennifer Widom slides

# Group-by versus subqueries

---

```
SELECT cName, count(*) AS cnt
FROM Apply
GROUP BY cName;
```

College( <u>cName</u> , state, enr)
Student( <u>sID</u> , sName, GPA, sizeHS)
Apply( <u>sID</u> , <u>cName</u> , <u>major</u> , decision)

```
SELECT DISTINCT cName,
    (SELECT count(*)
     FROM Apply A2
     WHERE A2.cName = A1.cName) AS cnt
FROM Apply A1;
```

cName	cnt
Berkeley	3
Cornell	6
MIT	4
Stanford	6

# Group-by versus subqueries

College(cName, state, enr)

Student(sID, sName, GPA, sizeHS)

Apply(sID, cName, major, decision)

SELECT cName

FROM Apply

GROUP BY cName

HAVING count(\*) < 5;



No need for DISTINCT: automatically  
from GROUP BY

cName
Berkeley
MIT

SELECT DISTINCT cName

FROM Apply A1

WHERE 5 > (SELECT count(\*) FROM Apply A2 WHERE A2.cName =  
A1.cName);

Every *group by* and *having* query can be written without using those  
clauses

# Group-by versus subqueries

College( <u>cName</u> , state, enr)
Student( <u>sID</u> , sName, GPA, sizeHS)
Apply( <u>sID</u> , <u>cName</u> , <u>major</u> , decision)

```
SELECT cName
```

```
FROM Apply
```

```
GROUP BY cName
```

```
HAVING count(*) < 5;
```

```
SELECT DISTINCT cName
```

```
FROM Apply A1
```

```
WHERE 5 > (SELECT count(*)  
           FROM Apply A2  
           WHERE A2.cName = A1.cName);
```

This is SQL by an expert

This is SQL by a novice

## Which way is more efficient?

How many times do we do a SFW query in each case?

With GROUP BY can be much more efficient!

# Aggregation summary

---

SELECT	S	→	attributes $a_1, \dots, a_k$ and/or aggregates over other attributes
FROM	$R_1, \dots, R_n$		
WHERE	$C_1$	→	any condition on the attributes in $R_1, \dots, R_n$
GROUP BY	$a_1, \dots, a_k$		
HAVING	$C_2$	→	any condition on the aggregate expressions

## Evaluation steps

1. Evaluate FROM-WHERE: apply condition  $C_1$  on the attributes in  $R_1, \dots, R_n$
2. GROUP BY the attributes  $a_1, \dots, a_k$
3. Apply condition  $C_2$  to each group
4. Compute aggregates in S and return the result

# Agenda

---

Introduction

The JOIN family of operators

Basic SQL Statement

Aggregation

Table Variables and Set Operators

Null values

Subqueries in WHERE clauses

Data Modification statements

Subqueries in FROM and SELECT clauses

# NULL values

---

Unless specified otherwise, any value in an attribute can take on the special value NULL

NULL usually means that the value is undefined or unknown or not applicable

We will see what happens when we have NULL values and we run queries over the database

# NULL values for numerical operations

---

NULL -> NULL

If  $x = \text{NULL}$  then

$4 \cdot (3 - x) / 7$  is NULL



# NULL values for Boolean operations

---

Conditions are evaluated using a three value logic:  
TRUE, FALSE, UNKNOWN

If  $x = \text{NULL}$  then  $x = \text{'Joe'}$  is UNKNOWN

Considering TRUE is 1.0, UNKNOWN is 0.5 and FALSE is 0.0

$C1 \text{ AND } C2 = \min(C1, C2)$

$C1 \text{ OR } C2 = \max(C1, C2)$

$\text{NOT } C1 = 1 - C1$

Rule in SQL: include only tuples that yield TRUE (1.0)

# A first query with NULL values

College(cName, state, enr)

Student(sID, sName, GPA, sizeHS)

Apply(sID, cName, major, decision)

SELECT sID, sName, GPA

FROM Student

WHERE GPA > 3.5;



sID	sName	GPA
123	Amy	3.9
234	Bob	3.6
456	Doris	3.9
678	Fay	3.8
987	Helen	3.7
876	Irene	3.9
654	Amy	3.9

2 additional students

Student

sID	sName	GPA	HS
123	Amy	3.9	1000
234	Bob	3.6	1500
345	Craig	3.5	500
456	Doris	3.9	1000
567	Edward	2.9	2000
678	Fay	3.8	200
789	Gary	3.4	800
987	Helen	3.7	800
876	Irene	3.9	400
765	Jay	2.9	1500
654	Amy	3.9	1000
543	Craig	3.4	2000
432	Kevin	NULL	1500
321	Lori	NULL	2500

## A second query with NULL values

---

```
SELECT sID, sName, GPA
FROM Student
WHERE GPA <= 3.5;
```

College(cName, state, enr)  
Student(sID, sName, GPA, sizeHS)  
Apply(sID, cName, major, decision)

<u>sID</u>	sName	GPA
345	Craig	3.5
567	Edward	2.9
789	Gary	3.4
765	Jay	2.9
543	Craig	3.4

## A third query with NULL values

College(cName, state, enr)

Student(sID, sName, GPA, sizeHS)

Apply(sID, cName, major, decision)

```
SELECT sID, sName, GPA
FROM Student
WHERE GPA > 3.5 OR GPA <= 3.5;
```



Tautology

Even with a tautology, we might not get all the data

<u>sID</u>	sName	GPA
123	Amy	3.9
234	Bob	3.6
345	Craig	3.5
456	Doris	3.9
567	Edward	2.9
678	Fay	3.8
789	Gary	3.4
987	Helen	3.7
876	Irene	3.9
765	Jay	2.9
654	Amy	3.9
543	Craig	3.4

College(cName, state, enr)

Student(sID, sName, GPA, sizeHS)

Apply(sID, cName, major, decision)

## A fourth query with NULL values

To make this query return all the students

```
SELECT sID, sName, GPA
```

```
FROM Student
```

```
WHERE GPA > 3.5 OR GPA <= 3.5 OR  
GPA IS NULL;
```

Can test for NULL explicitly

```
attribute IS NULL
```

```
attribute IS NOT NULL
```

sID	sName	GPA
123	Amy	3.9
234	Bob	3.6
345	Craig	3.5
456	Doris	3.9
567	Edward	2.9
678	Fay	3.8
789	Gary	3.4
987	Helen	3.7
876	Irene	3.9
765	Jay	2.9
654	Amy	3.9
543	Craig	3.4
432	Kevin	
321	Lori	

## A fifth query with NULL values

---

```
SELECT sID, sName, GPA, HS
FROM Student
WHERE GPA > 3.5 OR HS < 1600;
```

Although Kevin has a NULL GPA, he is retrieved because it has a HS<1600

sID	sName	GPA	HS
123	Amy	3.9	1000
234	Bob	3.6	1500
345	Craig	3.5	500
456	Doris	3.9	1000
678	Fay	3.8	200
789	Gary	3.4	800
987	Helen	3.7	800
876	Irene	3.9	400
765	Jay	2.9	1500
654	Amy	3.9	1000
432	Kevin		1500

College(cName, state, enr)

Student(sID, sName, GPA, sizeHS)

Apply(sID, cName, major, decision)

# NULL values and aggregate functions

---

SELECT distinct GPA  
FROM Student;

GPA
NULL
2.9
3.4
3.5
3.6
3.7
3.8
3.9

8 tuples

SELECT count(distinct GPA)  
FROM Student;

count(distinct GPA)
7

When counting (the distinct)  
values, NULLs are not included

College(cName, state, enr)  
Student(sID, sName, GPA, sizeHS)  
Apply(sID, cName, major, decision)

# NULL values

---

When using a database with NULL values

- Be careful when writing queries

- Understand how the NULL values are going to influence the result



# Agenda

---

Introduction

The JOIN family of operators

Basic SQL Statement

Aggregation

Table Variables and Set Operators

Null values

Subqueries in WHERE clauses

Data Modification statements

Subqueries in FROM and SELECT clauses

# Inserting new data

---

Insert Into **Table** ( $A_1, A_2, \dots, A_n$ )

Values ( $v_1, v_2, \dots, v_n$ )



If we provide values for all attributes, we may omit the list of attributes.

Insert Into **Table**

**Select-Statement**



Same schema as the table

# Inserting new data

---

Insert Into College

Values ('Carnegie Mellon', 'PA', 11500);

College(cName, state, enr)  
Student(sID, sName, GPA, sizeHS)  
Apply(sID, cName, major, decision)

College

cName	state	enr
Stanford	CA	15000
Berkeley	CA	36000
MIT	MA	10000
Cornell	NY	21000
Carnegie Mellon	PA	11500

# Inserting new data

College(cName, state, enr)  
Student(sID, sName, GPA, sizeHS)  
Apply(sID, cName, major, decision)

Have all students who didn't apply anywhere apply to CS at Carnegie Mellon

```
SELECT *  
  
FROM Student  
  
WHERE sID not in (select sID FROM Apply);
```

sID	sName	GPA	HS
456	Doris	3.9	1000
567	Edward	2.9	2000
789	Gary	3.4	800
654	Amy	3.9	1000

```
INSERT INTO Apply  
  
SELECT sID, 'Carnegie Mellon', 'CS', NULL  
  
FROM Student  
  
WHERE sID not in (select sID FROM Apply);
```

Apply

sID	cName	major	dec
...	...	...	...
543	MIT	CS	N
456	Carnegie Mellon	CS	NULL
567	Carnegie Mellon	CS	NULL
789	Carnegie Mellon	CS	NULL
654	Carnegie Mellon	CS	NULL

# Inserting new data

College(cName, state, enr)  
Student(sID, sName, GPA, sizeHS)  
Apply(sID, cName, major, decision)

Admit to Carnegie Mellon EE all students who were turned down in EE elsewhere

SELECT \*

FROM Student

WHERE sID in (select sID FROM Apply WHERE major='EE' AND  
decision='N' AND cName<>'Carnegie Mellon');

INSERT INTO Apply

SELECT sID, 'Carnegie Mellon', 'EE', 'Y'

FROM Student

WHERE sID in (select sID FROM Apply WHERE major='EE' AND  
decision='N' AND cName<>'Carnegie Mellon');

sID	sName	GPA	HS
123	Amy	3.9	1000
345	Craig	3.5	500

Apply

sID	cName	major	dec
...	...	...	...
654	Carnegie Mellon	CS	NULL
123	Carnegie Mellon	EE	Y
345	Carnegie Mellon	EE	Y

# Deleting existing data

---

Delete From **Table**

Where **Condition**

Condition can be complicated

Can include subqueries and aggregation over other tables

# Deleting existing data

College(cName, state, enr)  
Student(sID, sName, GPA, sizeHS)  
Apply(sID, cName, major, decision)

Delete all students who applied to more than two different majors

SELECT sID

FROM Apply

GROUP BY sID

HAVING count(distinct major) > 2;

DELETE FROM Student

WHERE sID in (

SELECT sID FROM Apply GROUP BY sID

HAVING count(distinct major) > 2);

**sID**

345

876

Student

sID	sName	GPA	HS
123	Amy	3.9	1000
234	Bob	3.6	1500
456	Doris	3.9	1000
567	Edward	2.9	2000
678	Fay	3.8	200
789	Gary	3.4	800
987	Helen	3.7	800
765	Jay	2.9	1500
654	Amy	3.9	1000
543	Craig	3.4	2000

# Deleting existing data

---

College(cName, state, enr)

Student(sID, sName, GPA, sizeHS)

Apply(sID, cName, major, decision)

Delete those students also from Apply

```
DELETE FROM Apply
```

```
WHERE sID in (
```

```
    SELECT sID FROM Apply GROUP BY sID
```

```
    HAVING count(distinct major) > 2);
```

Not all database systems allow deletion commands where the subquery includes the same relation that you're deleting from

In systems that don't allow, a temporary table would have to be created



# Deleting existing data

College(cName, state, enr)  
Student(sID, sName, GPA, sizeHS)  
Apply(sID, cName, major, decision)

Delete colleges with no CS applicants

SELECT \*

FROM College

WHERE cName not in

(SELECT cName from Apply where major='CS');

cName	state	enr
Cornell	NY	21000

DELETE FROM College

WHERE cName not in

(SELECT cName from Apply where major='CS');

cName	state	enr
Stanford	CA	15000
Berkeley	CA	36000
MIT	MA	10000
Carnegie Mellon	PA	11500

College

# Updating existing data

---

Update **Table**

Set **Attr = Expression**

Where **Condition**

Update **Table**

Set  **$A_1 = \text{Expr}_1, A_2 = \text{Expr}_2, \dots, A_n = \text{Expr}_n$**

Where **Condition**

Conditions and expressions can include subqueries and queries over other tables or the same table

# Updating existing data

College(cName, state, enr)  
Student(sID, sName, GPA, sizeHS)  
Apply(sID, cName, major, decision)

Accept applicants to Carnegie Mellon with GPA<3.6 but turn them into economics majors

```
SELECT *  
FROM Apply  
WHERE cName = 'Carnegie Mellon' and sID in  
(SELECT sID from Student where GPA<3.6);
```

```
UPDATE Apply  
SET decision = 'Y', major='economics'  
WHERE cName = 'Carnegie Mellon' and sID in  
(SELECT sID from Student where GPA<3.6);
```

sID	cName	major	dec
567	Carnegie Mellon	CS	
789	Carnegie Mellon	CS	

Apply			
sID	cName	major	dec
...	...	...	...
654	Carnegie Mellon	CS	NULL
123	Carnegie Mellon	EE	Y
567	Carnegie Mellon	economics	Y
789	Carnegie Mellon	economics	Y

# Updating existing data

---

Turn the highest GPA EE applicant into a CSE applicant

SELECT \*

FROM Apply

WHERE major='EE' AND sID in

(select sID from Student where GPA>=all

(select GPA from Student where sID in

(select sID from Apply where major='EE')));

sID	cName	major	dec
123	Stanford	EE	N
123	Cornell	EE	Y
123	Carnegie Mellon	EE	Y

College(cName, state, enr)

Student(sID, sName, GPA, sizeHS)

Apply(sID, cName, major, decision)

# Updating existing data

---

Turn the highest GPA EE applicant into a CSE applicant

UPDATE Apply

SET major = 'CSE'

WHERE major='EE' AND sID in

(select sID from Student where GPA>=all

(select GPA from Student where sID in

(select sID from Apply where major='EE')));

sID	cName	major	dec
...	...	...	...
789	Carnegie Mellon	economics	Y
123	Stanford	CSE	N
123	Cornell	CSE	Y
123	Carnegie Mellon	CSE	Y

Apply

College(cName, state, enr)

Student(sID, sName, GPA, sizeHS)

Apply(sID, cName, major, decision)

# Updating existing data

College(cName, state, enr)

Student(sID, sName, GPA, sizeHS)

Apply(sID, cName, major, decision)

Give every student the highest GPA and the smallest high school in the database

UPDATE Student

SET GPA = (select max(GPA) from Student),

sizeHS = (select min(sizeHS) from Student);

In the SET command, the right-hand side of the equals can itself be a subquery

Student

sID	sName	GPA	HS
123	Amy	3.9	200
234	Bob	3.9	200
456	Doris	3.9	200
567	Edward	3.9	200
678	Fay	3.9	200
789	Gary	3.9	200
987	Helen	3.9	200
765	Jay	3.9	200
654	Amy	3.9	200
543	Craig	3.9	200

# Summary

---

SQL is a rich programming language that handles the way data is processed declaratively

# Kahoot time!

---

Any doubts?



# Readings

---

Jeffrey Ullman, Jennifer Widom, A first course in Database Systems 3<sup>rd</sup> Edition

Section 6.1 – Simple Queries in SQL

Section 6.2 – Queries Involving More Than One Relation

Section 6.3 - Subqueries

Section 6.4 – Full-Relation Operations

Section 6.5 – Database Modifications

Philip Greenspun, SQL for Web Nerds,  
<http://philip.greenspun.com/sql/>