

SQL – Indexes

Carla Teixeira Lopes

Bases de Dados

Mestrado Integrado em Engenharia Informática e Computação, FEUP

Based on Jennifer Widom slides

Indexes

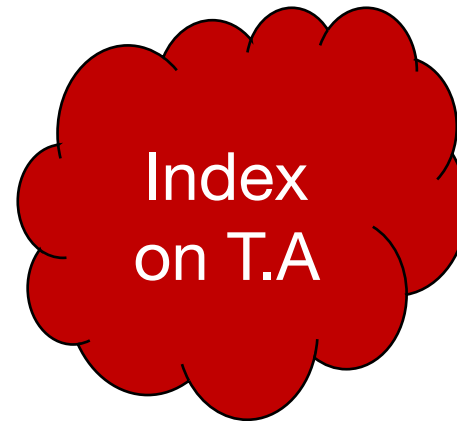
Primary mechanism to get improved performance on a database

Persistent data structure, stored in database

Many interesting implementation issues

But we are focusing on user/application perspective

Functionality



If we want tuples with T.A = 'cow',
DBMS doesn't need to scan the entire
table


Users don't access indexes

Indexes are used underneath by the
query execution engine

T

	A	B	C
1	cat	2	...
2	dog	5	...
3	cow	1	...
4	dog	9	...
5	cat	2	...
6	cat	8	...
7	cow	6	...

Functionality



Useful for queries involving T.B like:

$T.B = 2$

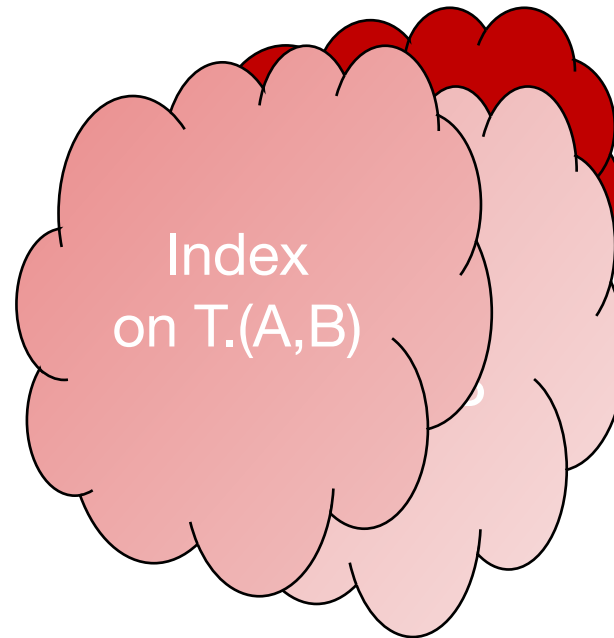
$T.B < 6$

$4 < T.B \leq 8$

T

	A	B	C
1	cat	2	...
2	dog	5	...
3	cow	1	...
4	dog	9	...
5	cat	2	...
6	cat	8	...
7	cow	6	...

Functionality



T

	A	B	C
1	cat	2	...
2	dog	5	...
3	cow	1	...
4	dog	9	...
5	cat	2	...
6	cat	8	...
7	cow	6	...

Useful for queries involving T.A and T.B
like:

T.A = 'cat' AND T.B > 5

T.A < 'd' AND T.B = 1

Utility

Index = difference between full table scans and immediate location of tuples

Orders of magnitude performance difference

Underlying data structures

Balanced trees (B trees, B+ trees)

For equalities or inequalities conditions

Operations running time tend to be logarithmic

Hash tables

Only for equality conditions

Operations running time is more or less constant

Where should indexes be created?

```
Select sName  
From Student  
Where sID = 18942
```

Many DBMS's build indexes automatically on PRIMARY KEY (and sometimes UNIQUE) attributes

Where should indexes be created?

```
Select sID  
From Student  
Where sName = 'Mary' and GPA>3.9
```

Index on sName

Hash or tree-based

Index on GPA

Tree-based

Index on (sName, GPA)

Where should indexes be created?

```
Select sName, cName  
From Student, Apply  
Where Student.sID = Apply.sID
```

More attributes →
Query planning &
optimization

Index on Apply.sID

Scans Student relation and, for each student, finds the matching SID in the Apply relation

Index on Student.sID

Scans Apply relation and, for each student, finds the matching SID in the Student relation

Index on (Student.sID, Apply.sID)

Downside of Indexes

Extra space

Marginal downside

Index creation

Medium downside

Index maintenance

Can offset benefits

Picking which indexes to create

Benefit of an index depends on:

Size of table (and possibly layout)

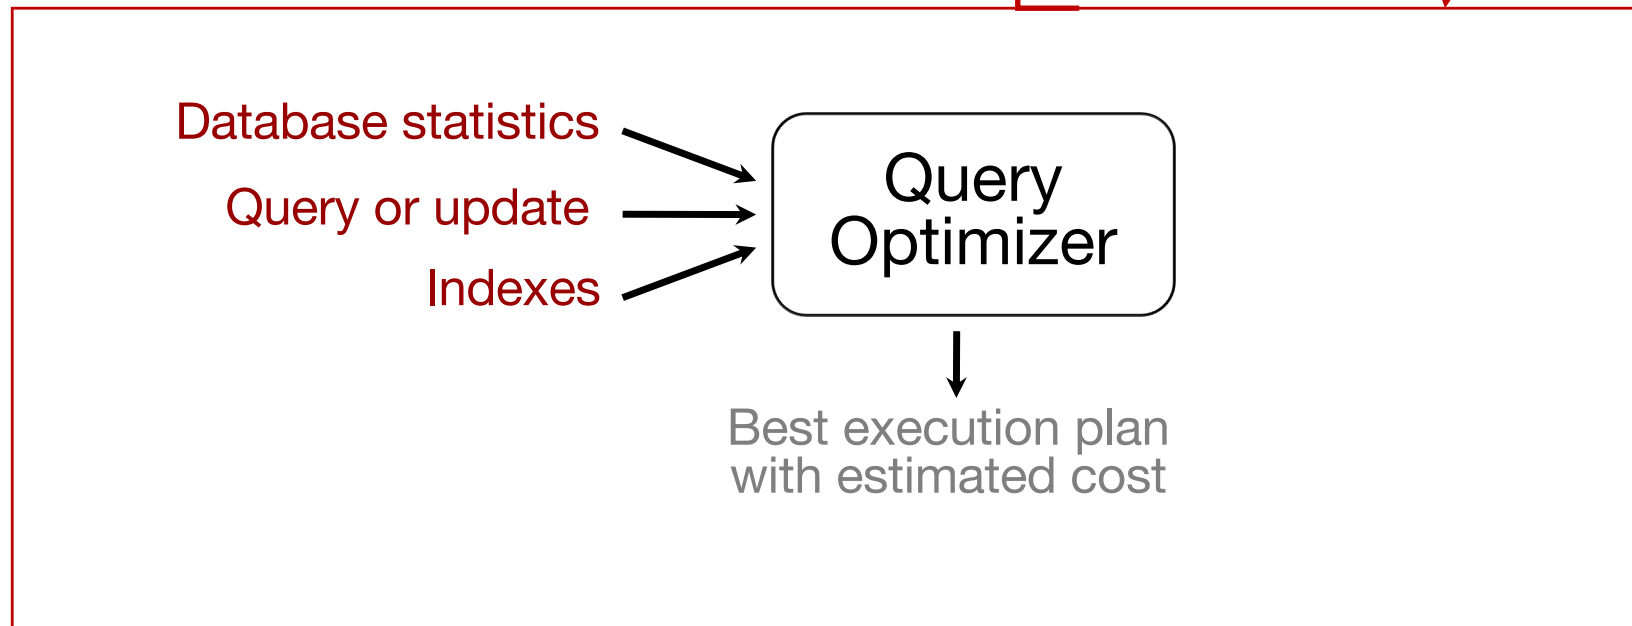
Data distributions

Query vs. update load

“Physical design advisors”

Input: database (statistics) and workload

Output: recommended indexes



Selects indexes whose
benefits outweigh
drawbacks

SQL Syntax

Create Index IndexName on T(A)

Create Index IndexName on T(A1,A2,...,An)

Create Unique Index IndexName on T(A)

Drop Index IndexName

Kahoot time!

Any doubts?

Readings

Jeffrey Ullman, Jennifer Widom, A first course in
Database Systems 3rd Edition

Section 8.3 – Indexes in SQL

Section 8.4 – Selection of Indexes